

COMPUTER VISION - ASSIGNMENT 2 - REPORT

2) MEAN SHIFT ALGORITHM

REMARK:

For all the functions, I have employed methods belonging to the **NUMPY** library. I have preferred to use NUMPY instead of TORCH as I have observed, after several attempts, that NUMPY performs better in terms of speed. Moreover, for the sake of consistency, I have used NUMPY both in the naive and the accelerated version. In that way, I can fairly compare the performances of the two versions.

2.1) In the distance function, I have calculated a vector which keep track of the distances between x and X , using the definition of euclidean norm. Using the `numpy.linalg.norm` function would have been more elegant, but I have observed that the manual implementation performs slightly better.

2.2) In the gaussian function, I have computed each weight according to a “Gaussian Kernel” with bandwidth 2.5. It is worth noticing that assigning the weights in this way penalizes points further away from x .

2.3) In the update function, I have updated x with the weighted averages of X . The weights used were of course those computed by the previous function. Once again, using functions like `numpy.dot` and `numpy.average` would have been smoother and more elegant, but I have observed that manual manipulation of the tensors performs better in terms of speed.

I have carefully avoided using for loops in the functions, as the performances steeply decline. In order to manually perform the distance and the update, I have often used a syntax like `weight[:,None]`, `X[None,:]` which is useful to add an extra dimension to the tensors.

2.4) In the accelerated version of the Mean-Shift algorithm, I had to redefine the distance and the update function as well as the `meanshift_step_batch`. In this function, I have defined the constants `BATCH_SIZE = 49`, `BATCH_NUMBER = 75`. Basically the mean shift algorithm now runs one batch at a time. It is worth noticing that the constants `BATCH_SIZE`, `BATCH_NUMBER` are not independent from each other. In fact, `BATCH_SIZE*BATCH_NUMBER = 3675` which equals the total number of pixels. I have run the script with several values of `BATCH_SIZE` and `BATCH_NUMBER`, finding out that 49 and 75 are optimal values in terms of speed.

After several attempts the average running time of both implementations are:

RUNNING TIME (CPU) of **NAIVE IMPLEMENTATION**: 12.9 / 13 seconds

RUNNING TIME (CPU) of **ACCELERATED IMPLEMENTATION**: 3.8/ 4 seconds

3) SEGNET

After having completed the model of the Segnet, I have trained and validated it. I have obtained a validation score of 0.871.